

# **SOCAL HOUSE PRICE ESTIMATOR**

**HYBRID MACHINE LEARNING MODEL  
TABULAR DATA + IMAGES**

**Javier García Esteve**



# PROJECT MOTIVATION

## Why Tabular Data Isn't Enough

- Price depends on numeric factors (bedrooms, bathrooms, sqft...)
- Location
- But also on visual factors, such as:
  - Condition & maintenance level
  - Architectural style
  - Landscaping & curb appeal
  - Overall quality

## Limitation of Tabular Models

- Tabular models capture structural attributes
- But they are blind to how the property actually looks

## Goal

- Combine both worlds:
  - Structured tabular data
  - Automatically extracted visual information from house images

→ A richer, more realistic home-price prediction model



# DATASET OVERVIEW

15,474 Southern California listings, each with:

- Tabular metadata:
- Bedrooms · Bathrooms · Sqft · City · Address · Price
- Target variable: price
- Property image

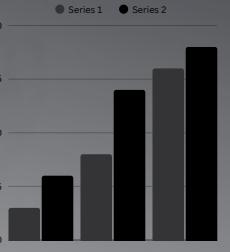
EDA insights:

- Price highly skewed → log transform
- City dominates price variation → need good encoding



# TABULAR MODELING (BASELINES)

Model	MAE	RMSE	R <sup>2</sup>
Linear Regression	0.3097	0.3981	0.4
Random Forest	0.3073	0.3905	0.4227
XGBoost	<b>0.2007</b>	<b>0.286</b>	<b>0.6904</b>



# FEATURE ENGINEERING & HYPERPARAMETER TUNING

Created features that reflect real estate logic:

- sqft\_per\_bed
- sqft\_per\_bath
- log\_sqft
- total\_rooms
- city\_target\_enc (mean log-price per city)

Feature importance shows:

- City encoding = strongest factor
- Size metrics = next strongest
- Spaciousness features add extra lift

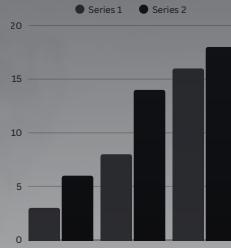
Hyperparameter Tuning

- Gridsearch

Result:

XGBoost + FE + tuning improved to  $R^2 = 0.78$

Model	MAE_test	RMSE_test	R2_test
<b>Linear Regression (baseline)</b>	0.309733	0.398098	0.400037
<b>Linear Regression (feat-eng)</b>	0.209571	0.286105	0.69012
<b>Random Forest (baseline)</b>	0.307297	0.390504	0.42271
<b>Random Forest (feat-eng)</b>	0.201842	0.27573	0.712187
<b>XGBoost (baseline)</b>	0.200721	0.285987	0.690375
<b>XGBoost (feat-eng)</b>	0.171125	0.24494	0.772876
<b>XGBoost (feat-eng, tuned)</b>	<b>0.166249</b>	<b>0.242512</b>	<b>0.777357</b>



# WHY INCORPORATE IMAGES?

Tabular data misses key visual factors:

- Renovation quality
- Architectural style
- Condition
- Landscaping
- Curb appeal
- Luxury finishes
- Neighborhood visual cues

Users often can't provide this information manually.  
Even homeowners may not know or can't objectively  
describe these attributes.

But everyone can take a picture!  
An image can capture many of these visual signals  
instantly.

Solution:  
let the model learn these visual cues automatically from  
images.



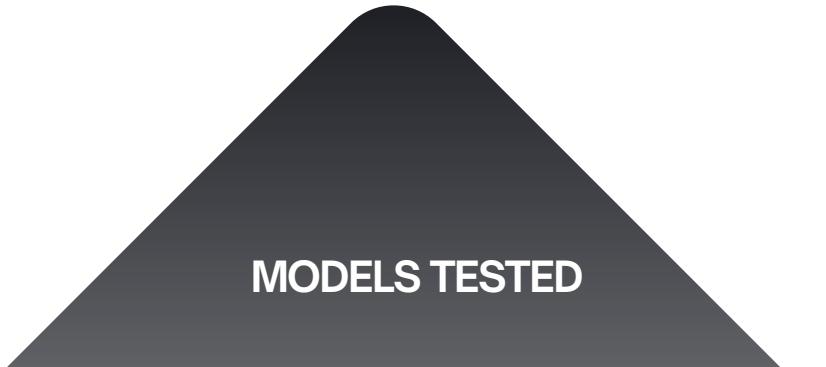
# CNN MODEL ARCHITECTURE

## Transfer Learning: EfficientNetB0 tuned

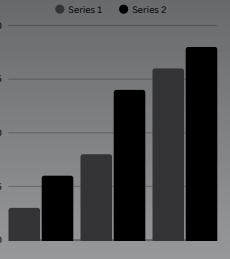
- EfficientNet backbone
- Unfreeze last 10 layers
- Lower learning rate
- Strong augmentations (flip, zoom, rotation, contrast)
- Added head:
  - GlobalAveragePooling
  - Dense layer(128) → embedding
  - Dropout 0.3
  - Dense layer(1) → log-price output

CNN models performed poorly for standalone price prediction.

The 128-dim embedding is the valuable output — used as visual features inside the hybrid model.



Model Texto	MAE_log_price	RMSE_log_price	R2_log_price
<b>EfficientNetB0 (baseline)</b>	0.712699	0.941124	-2.353031
<b>EfficientNetB3</b>	0.73529	0.985476	-2.676506
<b>EfficientNetB0 Tuned (last 10 unfrozen)</b>	<b>0.604645</b>	<b>0.781624</b>	<b>-1.312809</b>



# HYBRID MODEL

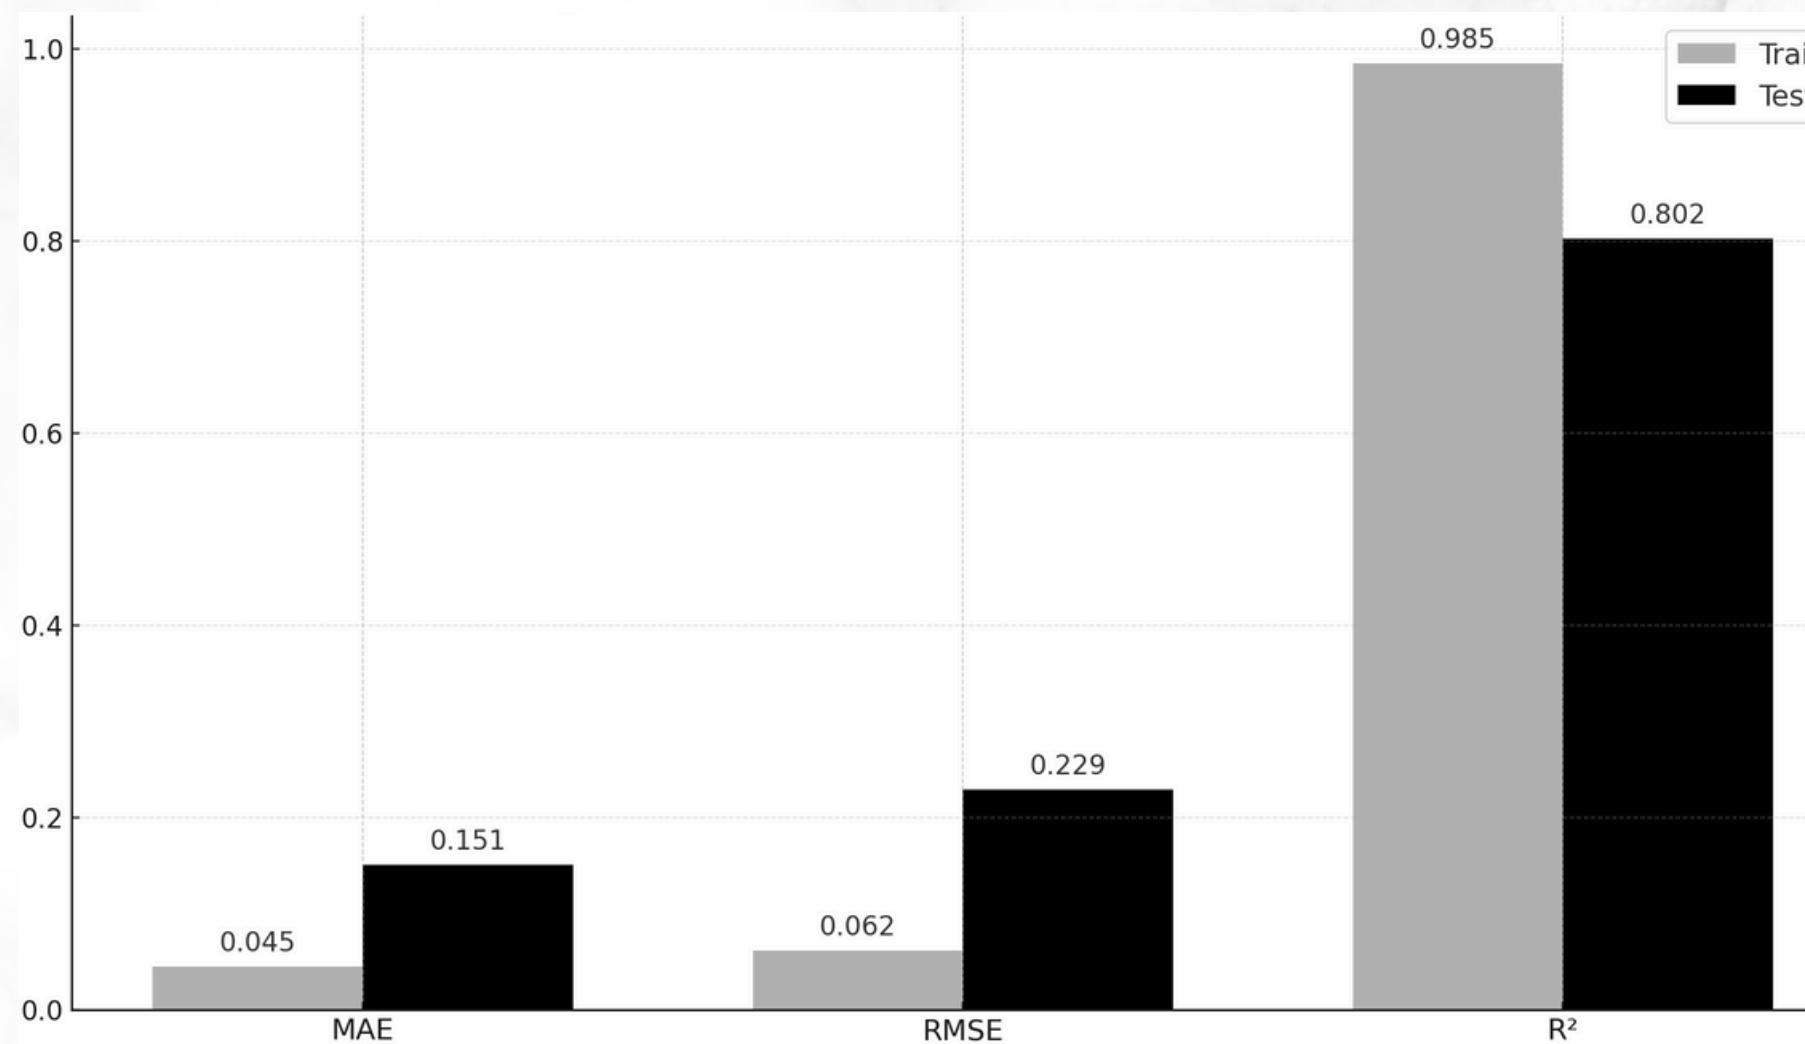
Hybrid = Tabular Features + Image  
Embedding → XGBoost

Steps:

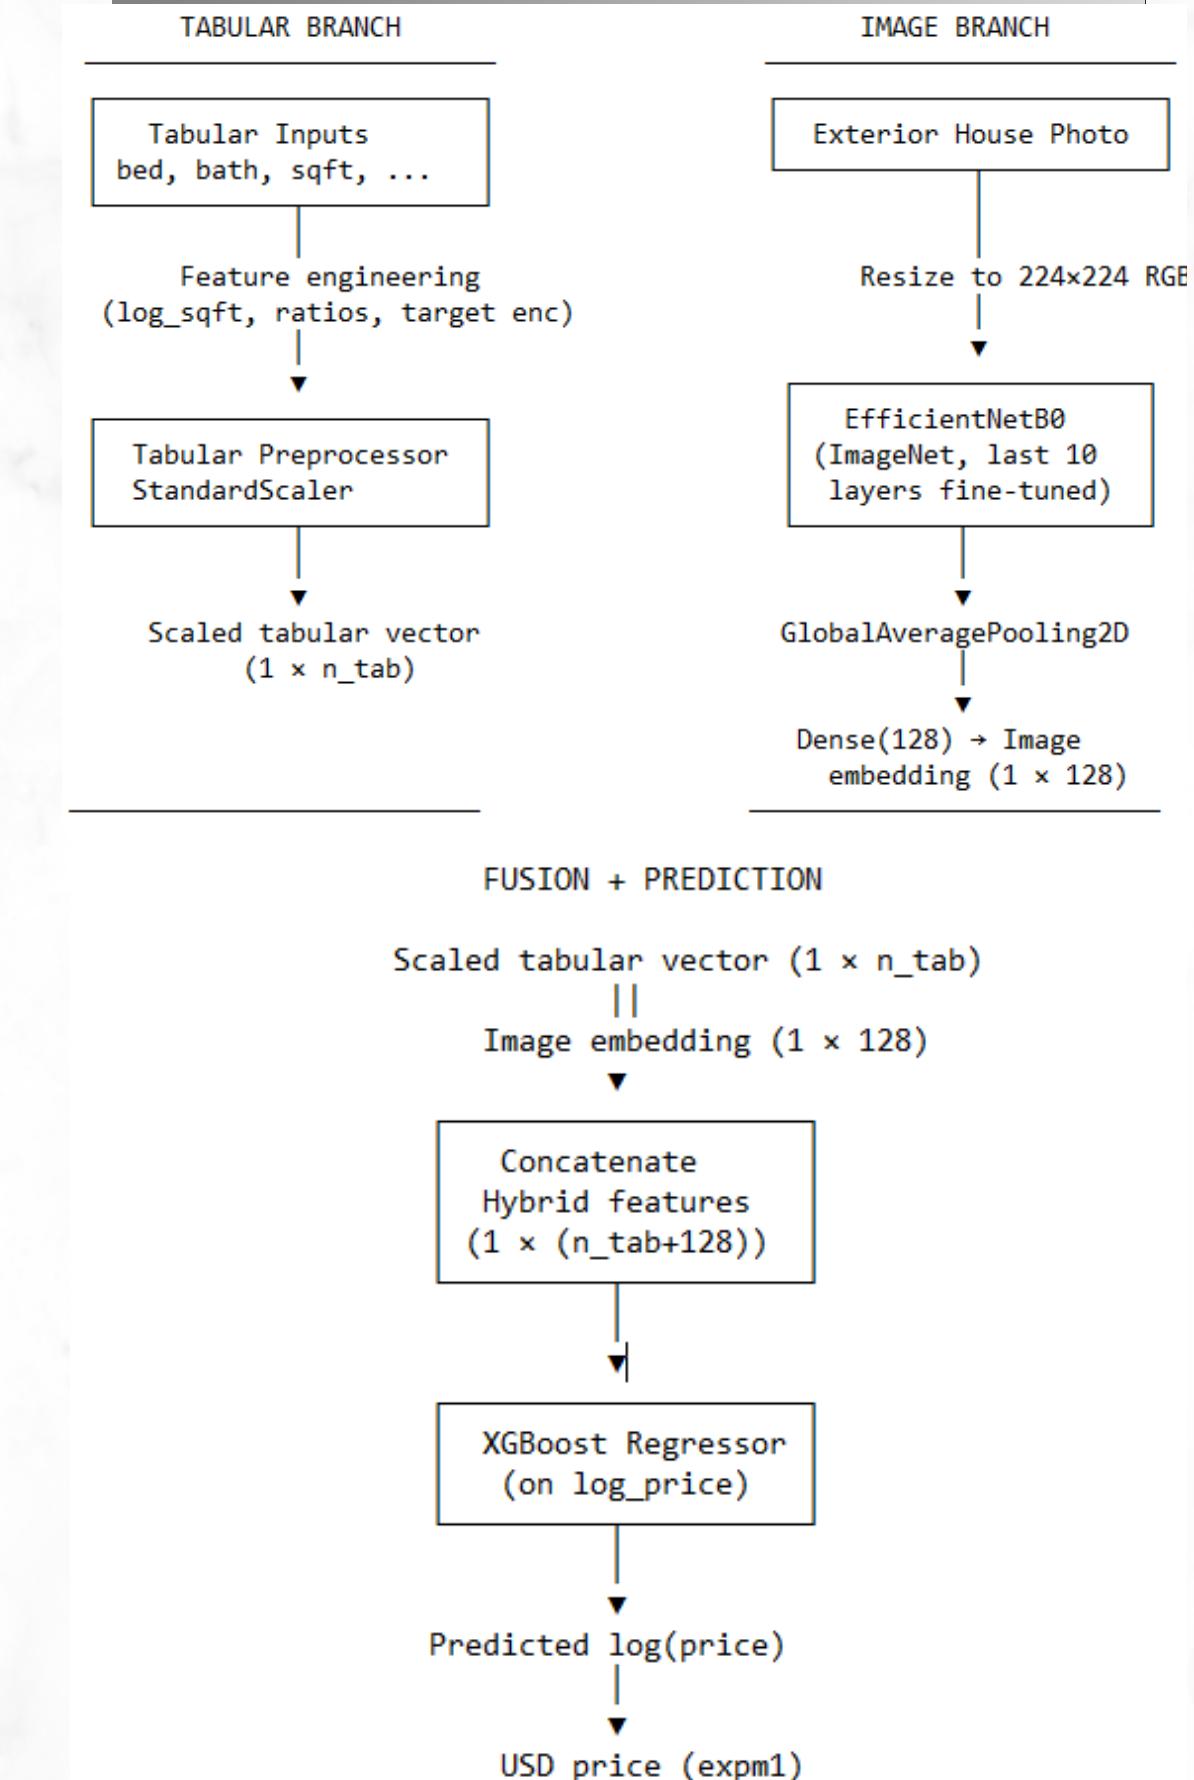
1. Scale tabular features
2. Extract 128-dim embedding from EfficientNetB0 (fine-tuned)
3. Combine tabular features + image embedding into one vector
4. Feed this hybrid vector into XGBoost regressor
5. Predict log(price) → convert back to USD

Performance:

- Hybrid Model:  $R^2 = 0.80$

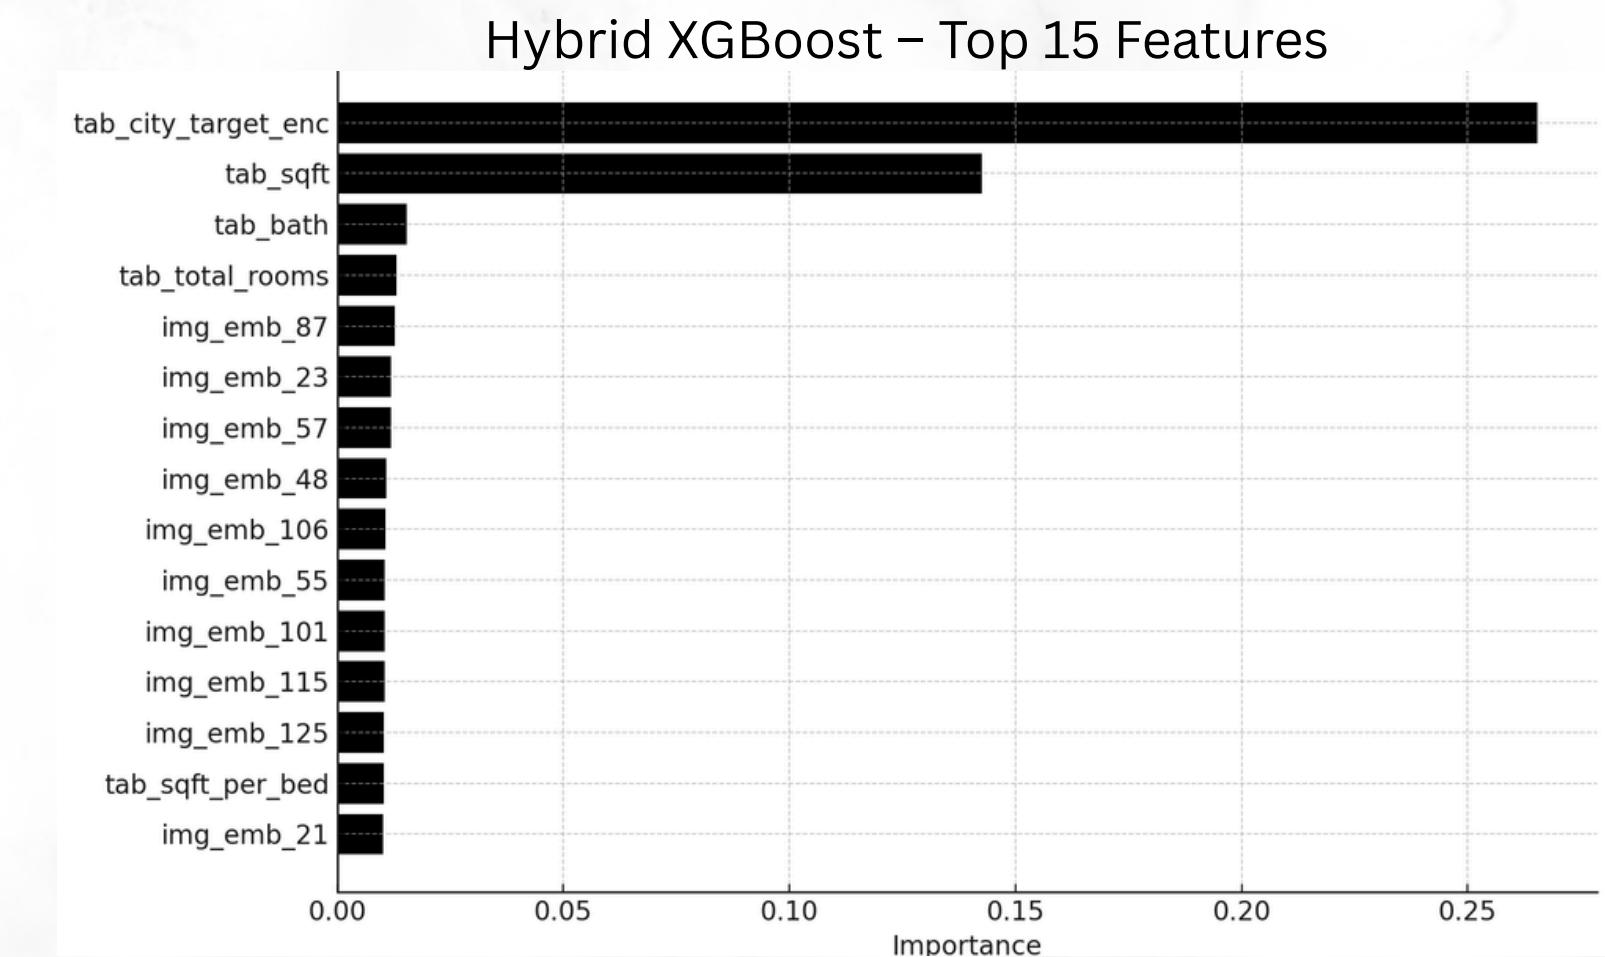
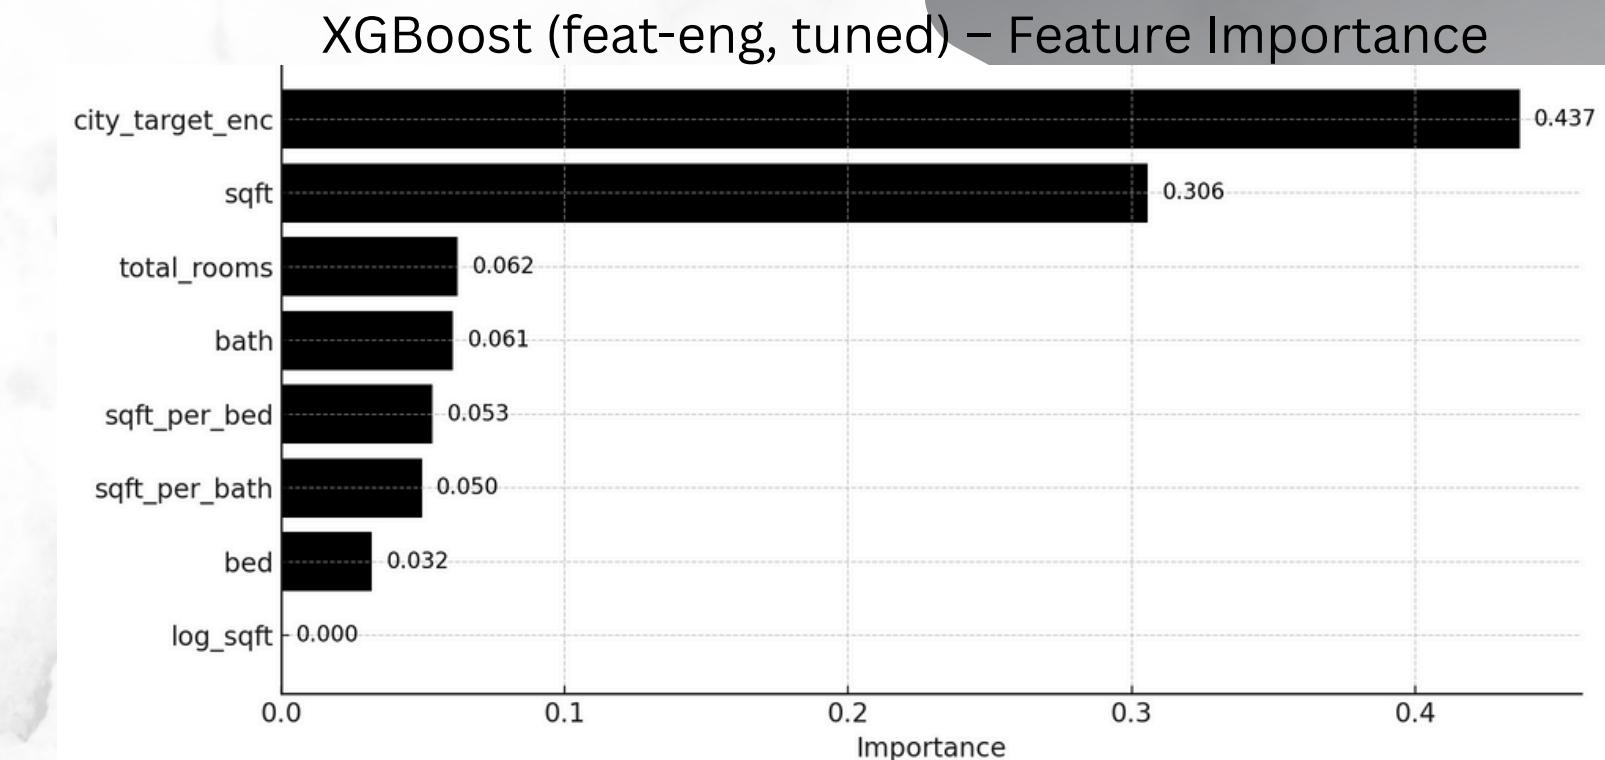


HYBRID ARCHITECTURE DIAGRAM



# TABULAR-ONLY VS HYBRID MODEL

Model	MAE_test	RMSE_test	R2_test
XGBoost (feat-eng, tuned - tabular only)	0.166249	0.242512	0.777357
Hybrid (tabular FE + image embeddings)	0.150762	0.228766	0.80188



# DEMO

“To make this accessible, I built a Streamlit app where the user inputs the property details and optionally uploads an image. If an image is provided you can choose to either run the hybrid model or just the tabular one. If you don’t provide an image, just the tabular-only model runs.”

The screenshot shows the SoCal House Price Estimator application. At the top, there's a logo of a house with a blue roof and red door, followed by the text "SoCal House Price Estimator". Below it, a subtitle reads "Hybrid ML model combining tabular features and pictures".

**1 Enter listing details**

This section contains input fields for property details:

- City:** A dropdown menu showing "29 Palms, CA".
- Living area (sqft):** An input field with "1500" and minus/plus buttons.
- Bedrooms:** An input field with "3" and minus/plus buttons.
- Bathrooms:** An input field with "2" and minus/plus buttons.
- Optional: upload a picture of the house:** A section with a placeholder "Drag and drop file here" and a "Browse files" button. It also specifies "Limit 200MB per file • JPG, JPEG, PNG".

---

**2 Get price estimate**

A checkbox labeled "Use hybrid model when a photo is uploaded" is checked. Below it is a large "Estimate price" button.

# CONCLUSION & FUTURE WORK

## Key Conclusions

- Feature engineering significantly boosts performance
- CNN embeddings add visual intelligence
- Hybrid model is the global best performer
- Real-time app demonstrates practical application



## Future Ideas

- Add Interior pictures
- Multiple images per listing
- Satellite imagery
- Multi-task models (price + condition scoring)

# THANK YOU!

